

Vaishnavi Khindkar

(412) 657-1340 • vkhindka@andrew.cmu.edu • vaishnvi.github.io • LinkedIn • Google Scholar • github.com/Vaishnvi

EDUCATION

Carnegie Mellon University, School of Computer Science (Robotics Institute)

Aug 2025 - Dec 2026

Master of Science in Computer Vision

GPA 4.11/4

- Relevant Coursework: Advanced Computer Vision (16-820), Learning for 3D Vision (16-825), Robot Learning (16-831), Robot Localization & Mapping (16-833), Visual Learning & Recognition (16-824), Deep Learning Systems (10-714).

TECHNICAL SKILLS

Programming: Python (3.x), Modern C++ (14, 17, 20), MATLAB, SQL, C, R, Octave, Core Java

ML / CV: Machine Learning, Deep Learning, Computer Vision, 3D Vision, Multimodal Learning, Object Detection

Frameworks and Tools: PyTorch, TensorFlow, AWS, NVIDIA Triton, Linux, Git, LaTeX, Jupyter Notebook, OpenCV, Docker

Domains: Generative Modeling & Sampling, 3D & 4D Scene Understanding, Video Generative AI, Object-Centric & Spatial Representations, World Models, Autonomous Driving, Robotics

EXPERIENCE

Sync Labs

San Francisco Bay Area, CA

Machine Learning Engineering Intern

May 2026 - Present

- Architected GPU-served VLM inference on Modal (serverless H100s): deployed multimodal models (Gemma-3-12B, Qwen2.5-VL-7B) for automated visual QC of lip-sync video, benchmarked against NVIDIA Cosmos Reason, with structured-output prompting, multi-model agreement gating, and per-model routing.
- Built a human-in-the-loop labeling loop and LoRA fine-tuning for Qwen2.5-VL on artifact crops; evaluated RLHF vs. SFT and built a per-class precision/recall/F1 evaluation harness.

Carnegie Mellon University, Machine Learning Department

Pittsburgh, PA

Graduate Research Fellow (Advisor: Prof. [Andrej Risteski](#))

May 2026 - Present

- Studying reward-tilted sampling: principled lift-and-tilt methods for steering generative models toward target distributions under complex, indefinite rewards.

Carnegie Mellon University, Robotics Institute

Pittsburgh, PA

Research Assistant (Advisor: Prof. [Fernando De la Torre](#))

Aug 2025 - Present

- Developed the first framework using foundation generative models (Imagen-3) for systematic diagnosis of aerial object detectors, enabling fine-grained failure analysis and up to 11% AP50 gains.

Product Labs (Bhashini, Government of India)

Hyderabad, India

Applied AI Researcher

Oct 2024 - Jul 2025

- Containerized OCR models with Docker and deployed via NVIDIA Triton, achieving real-time latency (0.11s printed, 0.051s handwritten) across 22 Indian languages.

CVIT, IIT Hyderabad

Hyderabad, India

Jr. Research Scientist (Advisors: Prof. [C. V. Jawahar](#), Prof. [Vineeth N. Balasubramanian](#), Prof. [Chetan Arora](#))

Aug 2020 - Sep 2024

- Led research on reason-enriched pedestrian intent estimation (MindRead, IROS 2024), creating the PIE++ dataset and achieving +5.6% accuracy and +7% F1 over baselines.
- Developed ILLUME, a self-attention-based adaptive object detection framework (WACV 2022) with improved structural attention and state-of-the-art domain adaptation performance.

Barclays

Pune, India

Software Developer

Aug 2018 - Oct 2020

- Built a Contacts application with Spring Boot and AngularJS, improving UX and reducing system errors.

SELECT PUBLICATIONS AND PATENTS

- Diagnosing Aerial-View Object Detectors with Foundational Image Generative Models [\[Link\]](#)
Panev S., Jeon M., Khindkar V., Deshpande A., de Melo C.M., Hu S., Chakraborty S., De la Torre F. (**ECCV 2026**)
- Can Reasons Help Improve Pedestrian Intent Estimation? A Cross-Modal Approach [\[Link\]](#)
Khindkar V., Balasubramanian V.N., Arora C., Subramanian A., Jawahar C.V. (**IROS 2024**)
- To miss-attend is to misalign! Residual Self-Attentive Feature Alignment for Adapting Object Detectors [\[Link\]](#)
Khindkar V., Balasubramanian V.N., Arora C., Saluja R., Subramanian A., Jawahar C.V. (**WACV 2022**)
- GAMMA: Generative Augmentation for Attentive Marine Debris Detection [\[Link\]](#)
Khindkar V., Khindkar J. (**arXiv 2022**)
- System and method for detecting object in an adaptive environment using a machine learning model
Khindkar V., Balasubramanian V.N., Arora C., Subramanian A., Jawahar C.V. (**US Patent 2023**)

SELECT PROJECTS

- **Slot4D: Learning Object Slots for 4D Reconstruction** - CMU Capstone (Advisor: Prof. [Laszlo A. Jeni](#))
Implementing an object-centric 4D reconstruction pipeline using slot attention to decompose dynamic scenes into per-object latent slots from multi-view video.
- **Cross-Resolution Gaussian Splatting for Hybrid Resolution Photography** - (Advisor: Prof. [Shubham Tulsiani](#))
Designed a cross-resolution 3D Gaussian Splatting pipeline that fuses heterogeneous image sources, incorporating detail-preserving densification for high-fidelity reconstruction.
- **Generalizable Single-View Holistic Robot Pose Estimation** - (Advisor: Prof. [Jun-Yan Zhu](#))
Designed a single-view robot pose and keypoint estimation pipeline with synthetic-to-real transfer and domain randomization for unseen robot types.

ACADEMIC SERVICE & VOLUNTEERING

Reviewer: CVPR 2025, ECCV 2024, CVPR 2023, ECCV 2022, WACV 2022 **Mentor:** AI4ALL Changemakers in AI (2022)

Teaching: CV tutorials for IISc-TalentSprint Exec. M.Tech **Organizer:** Hour of Code (Intl. Coding Week)